

# **Build and Manage your APIs with Amazon API Gateway**

# Agenda


Why we built Amazon API Gateway

What is Amazon API Gateway?

Amazon API Gateway Features & Functionality

Q&A

# At AWS, We run a lot of APIs



...Over time, we have learned a few lessons

# Your Feedback



Managing multiple versions and stages of an API is difficult



Monitoring 3<sup>rd</sup> party developers' access is time consuming



Access authorization is a challenge



Traffic spikes create operational burden



What if I don't want servers at all?

# Introducing Amazon API Gateway



- ✓ Host multiple versions and stages of your APIs
- ✓ Create and distribute API Keys to developers
- ✓ Leverage AWS Sigv4 to authorize access to APIs
- ✓ Throttle and monitor requests to protect your backend
- ✓ Utilizes AWS Lambda



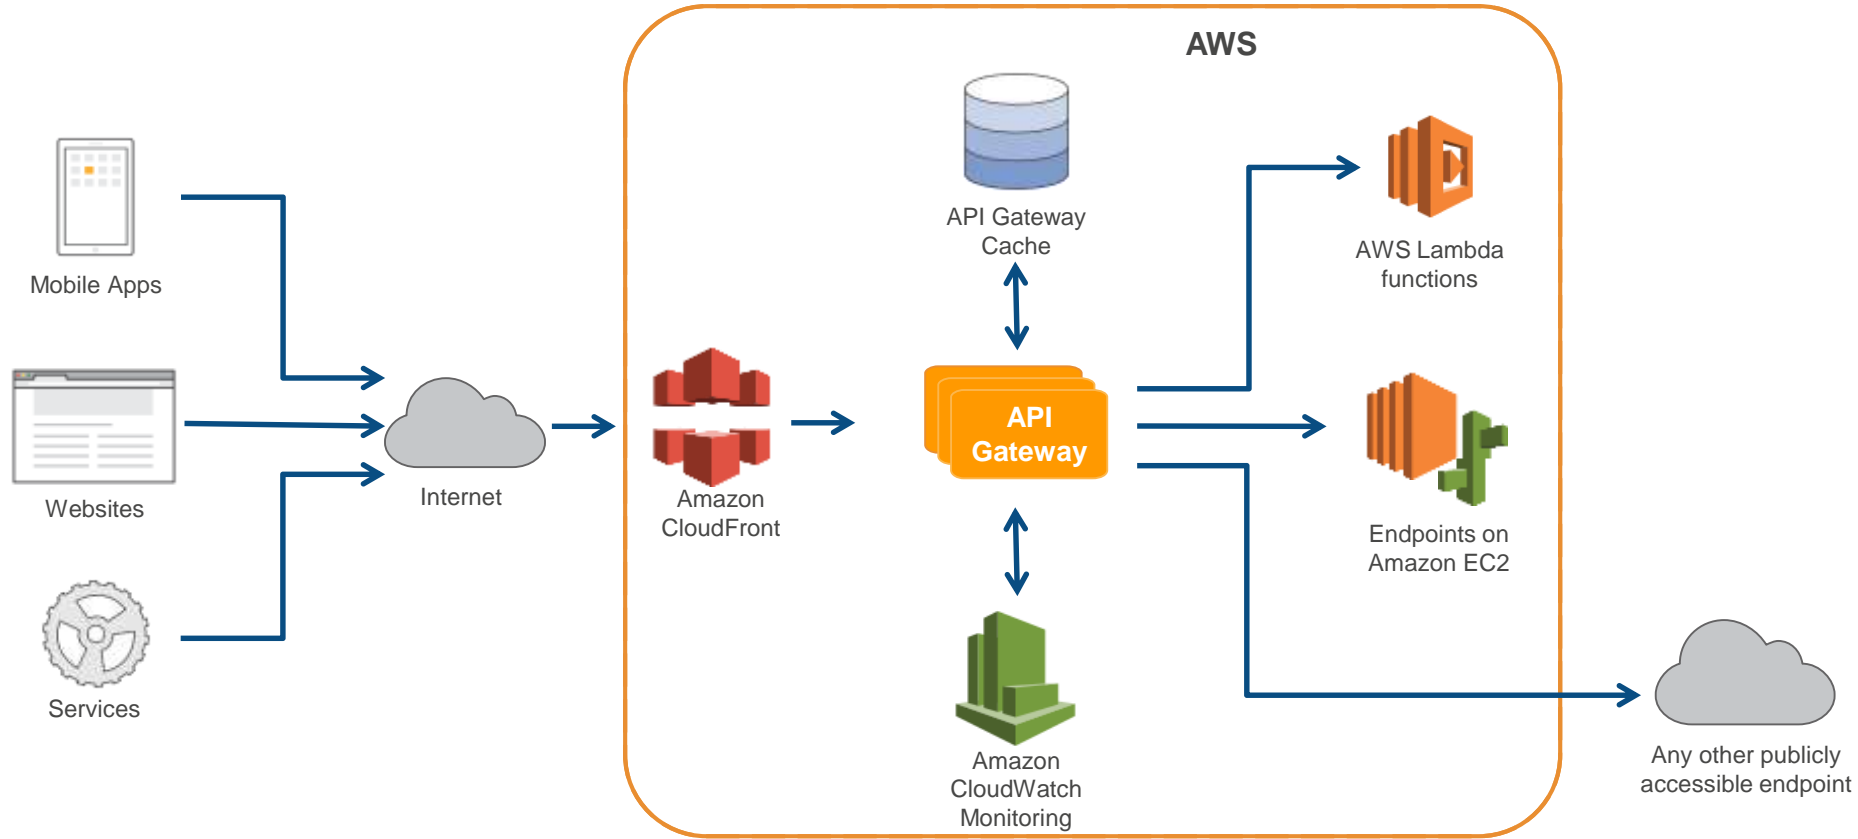
# Introducing Amazon API Gateway



- ✓ Managed cache to store API responses
- ✓ Reduced latency and DDoS protection through CloudFront
- ✓ SDK Generation for iOS, Android and JavaScript
- ✓ Swagger support
- ✓ Request / Response data transformation and API mocking

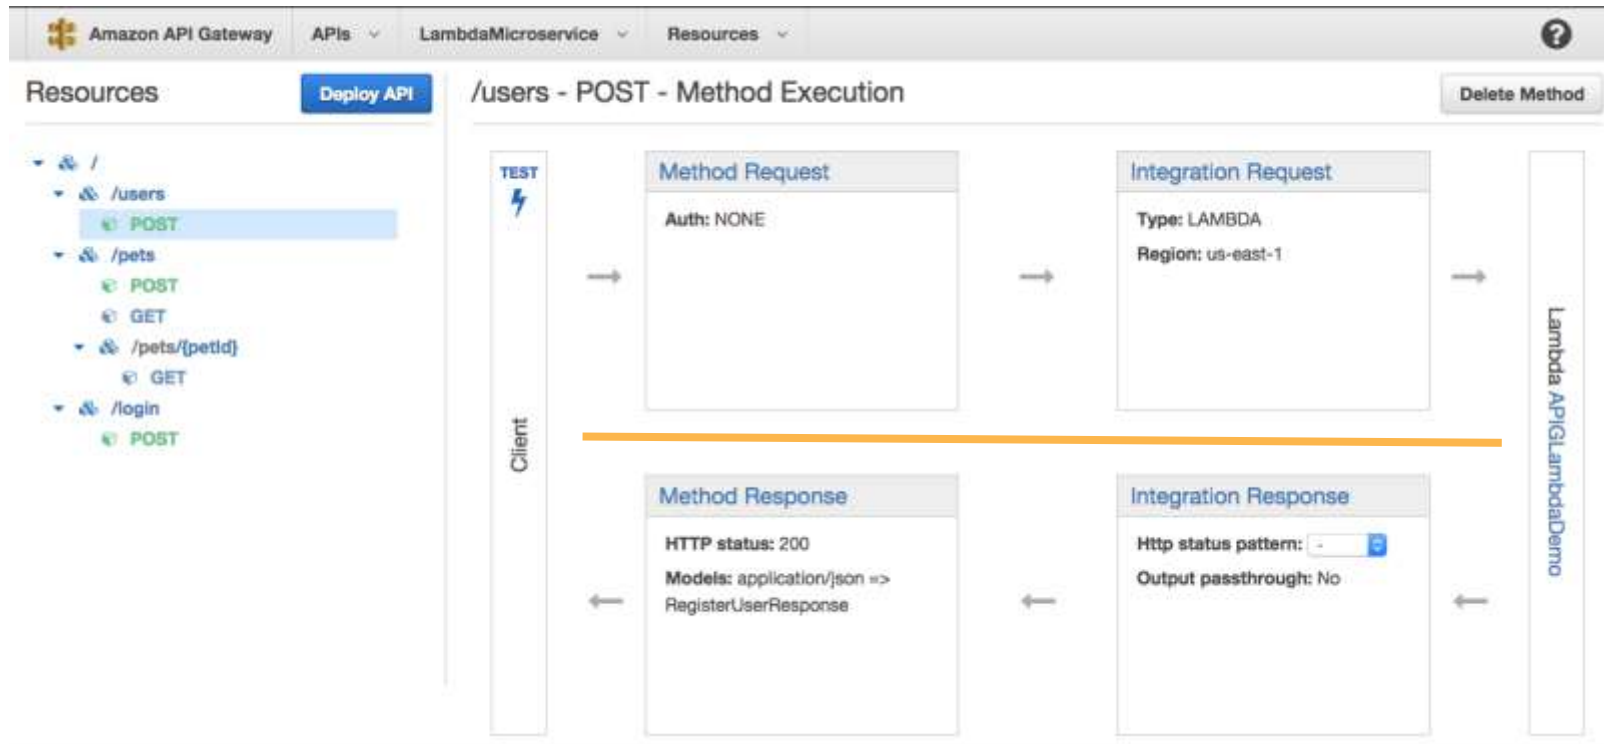
# How Does Amazon API Gateway Work?

# An API Call Flow

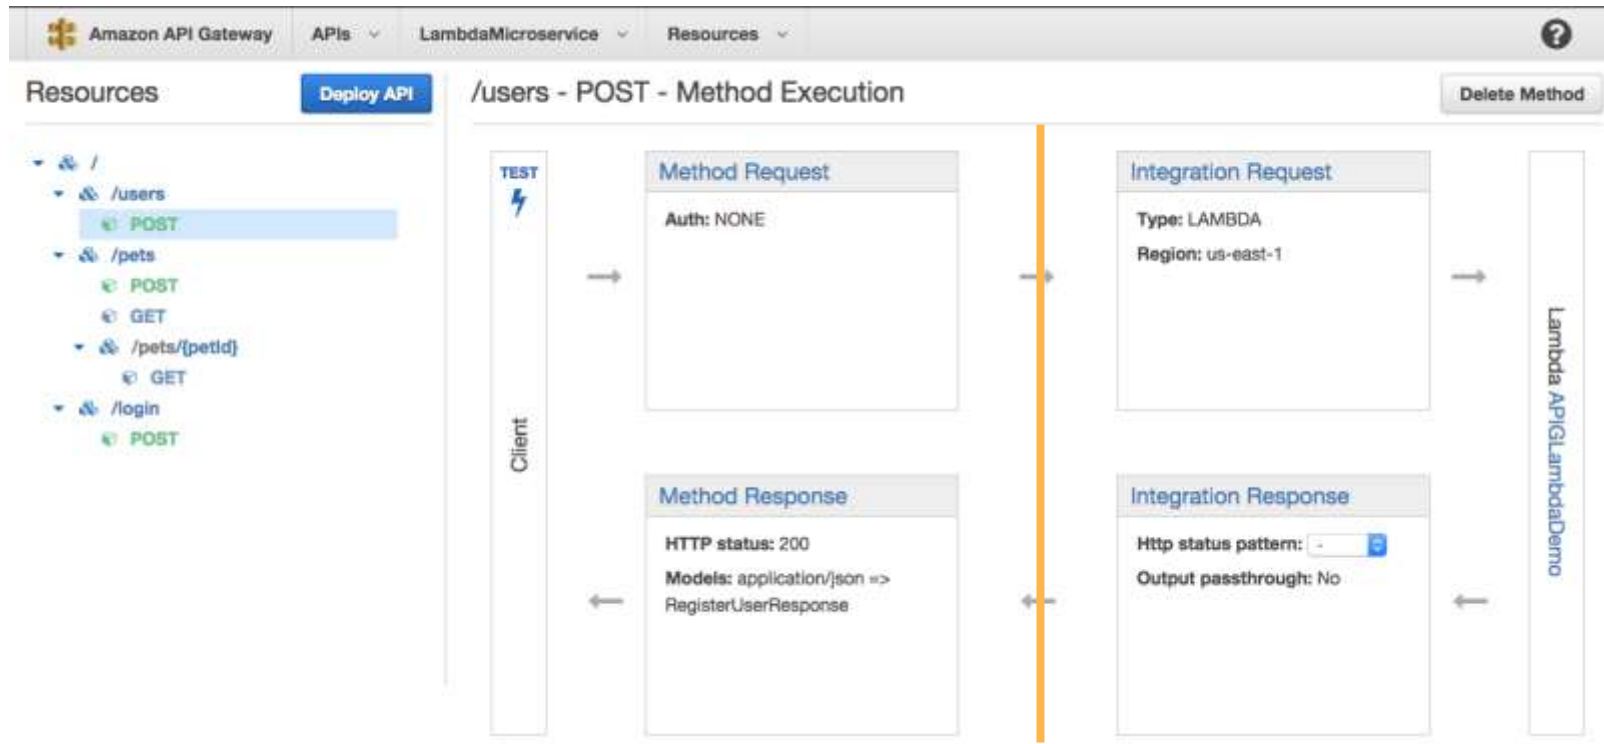




# Methods and Integrations



# Methods and Integrations



# Build, Deploy, Clone & Rollback

Build APIs with their resources, methods, and settings

Deploy APIs to a Stage

- Users can create as many Stages as they want, each with its own Throttling, Caching, Metering, and Logging configuration

Clone an existing API to create a new version

- Users can continue working on multiple versions of their APIs

Rollback to previous deployments

- We keep a history of customers' deployments so they can revert to a previous deployment

# API Configuration

You can create APIs

Define resources within an API

Define methods for a resource

- Methods are Resource + HTTP verb

Pet Store

/pets

/pets/{petId}

- **GET**
- **POST**
- **PUT**

# API Deployments

API Configuration can be deployed to a stage

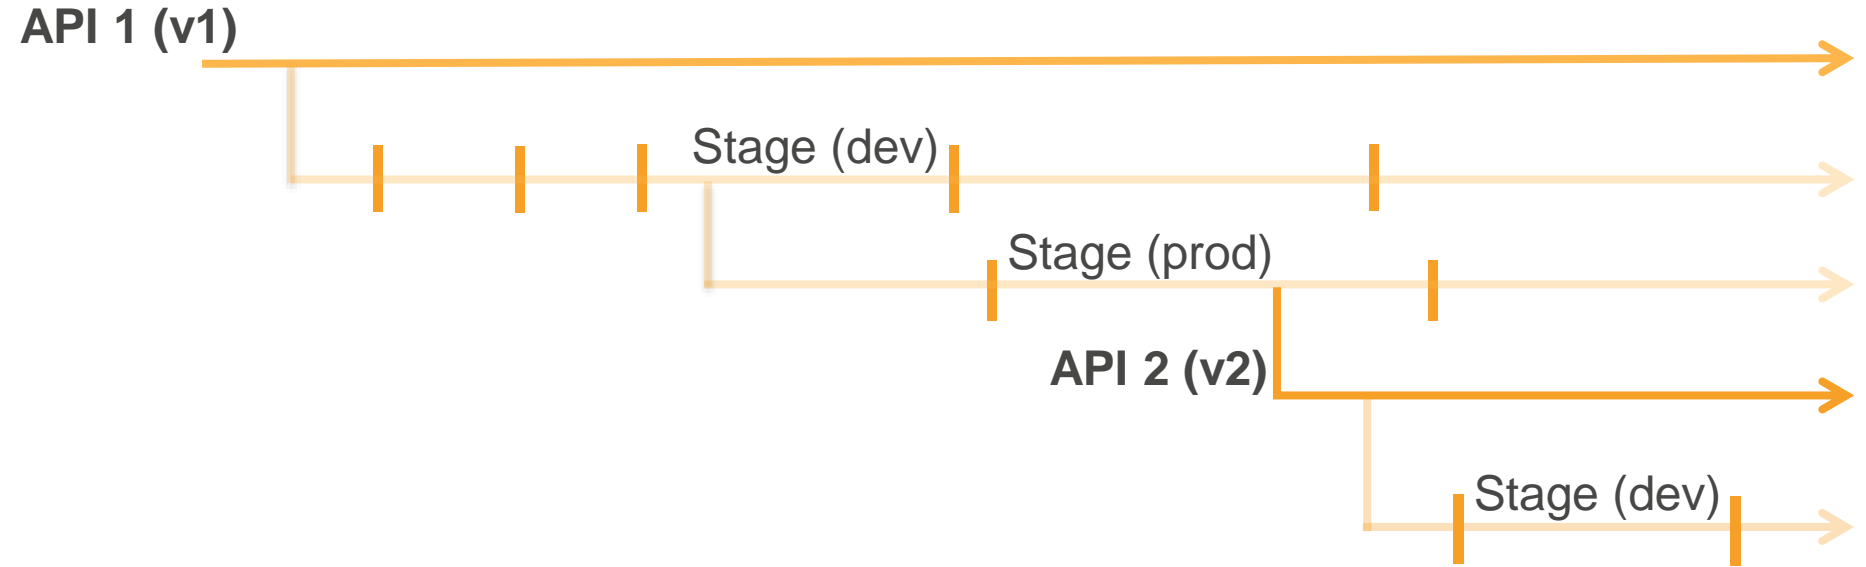
Stages are different environments

For example:

- Dev (e.g. [awsapigateway.com/dev](https://awsapigateway.com/dev))
- Beta (e.g. [awsapigateway.com/beta](https://awsapigateway.com/beta))
- Prod (e.g. [awsapigateway.com/prod](https://awsapigateway.com/prod))
- As many stages as you need



# Manage Multiple Versions and Stages of your APIs



# Custom Domain Names

You can configure custom domain names with subdomains and base paths

Pointing to an API you have access to all Stages

- Beta (e.g. yourapi.com/beta)
- Prod (e.g. yourapi.com/prod)

Pointing directly to your “prod” Stage

- Prod (e.g. yourapi.com/)

# Metering and Authorization



# API Keys to Meter Developer Usage

Create API Keys

Set access permissions at the API/Stage level

Meter usage of the API Keys through CloudWatch Logs

# API Keys



API Keys should be used purely to meter app/developer usage



API Keys should be used alongside a stronger authorization mechanism

# Leverage AWS Sigv4, or Use a Custom Header

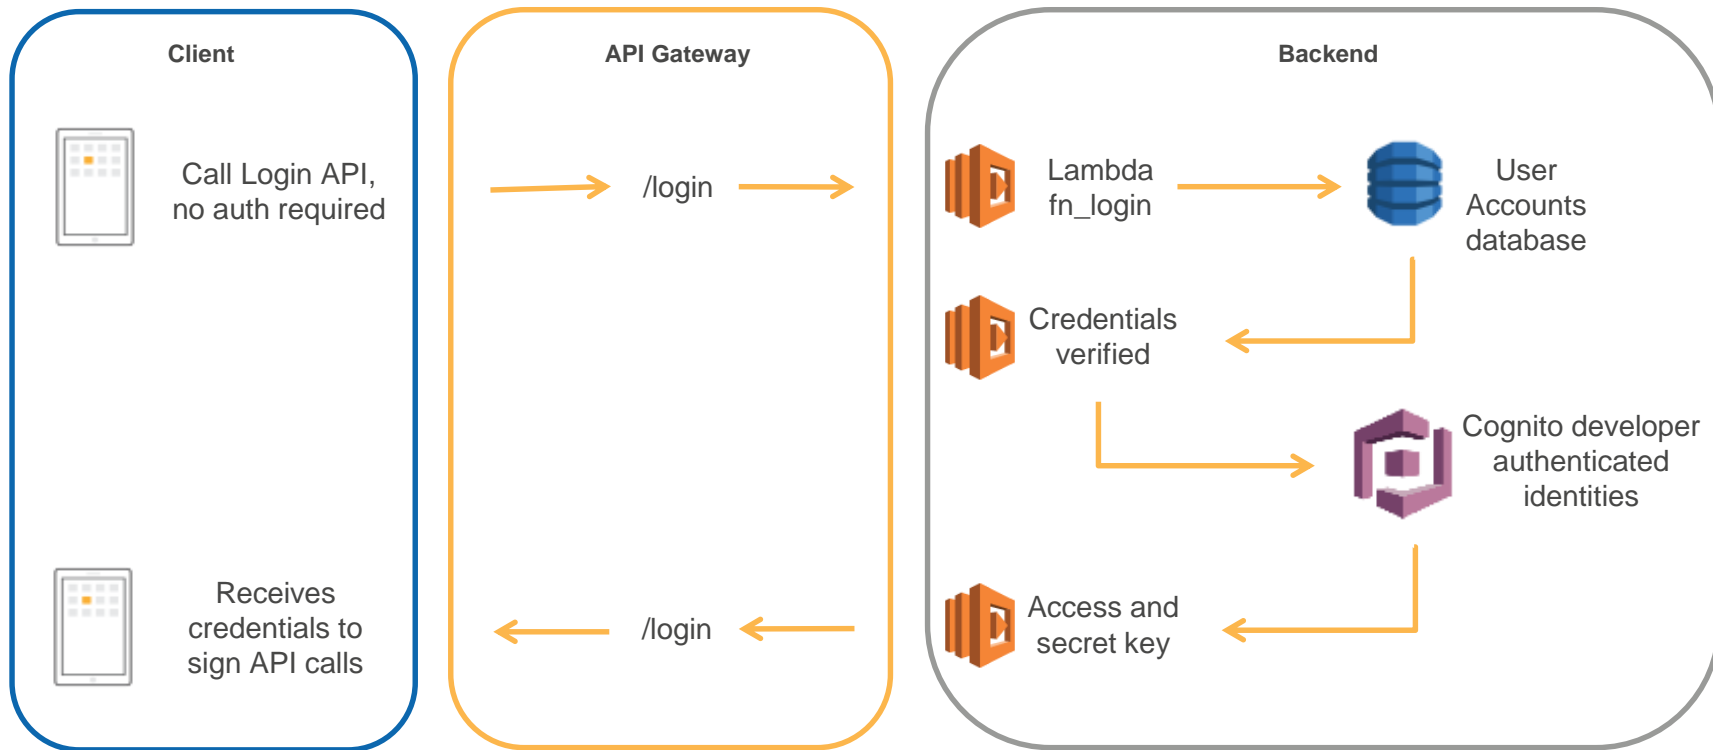
You can leverage AWS Sigv4 to sign and authorize API calls

- Amazon Cognito and AWS Security Token Service (STS) simplify the generation of temporary credentials for your app

You can support OAuth or other authorization mechanisms through custom headers

- Simply configure your API methods to forward the custom headers to you backend

# Using Sigv4 to authenticate calls to your API



# The AWSCredentialsProvider

We implement the AWSCredentialsProvider interface

```
@interface APISessionCredentialsProvider : NSObject <AWSCredentialsProvider>
```

The refresh() method is called whenever the SDK needs new credentials

```
- (AWSTask *)refresh {
    PETLambdaMicroserviceClient *client = [PETLambdaMicroserviceClient clientForKey:
        APIClientConfigurationKey];
    PETRegisterUserRequest *req = [PETRegisterUserRequest new];
    req.username = _credentials.username;
    req.password = _credentials.password;
    return [[client loginPost:req] continueWithBlock:^(id(AWSTask *)task) {
        PETLoginUserResponse *resp = task.result;

        PETLoginUserResponse_credentials *credentials = resp.credentials;

        _accessKey = credentials.accessKey;
        _secretKey = credentials.secretKey;
        _sessionKey = credentials.sessionToken;
        _expiration = [NSDate dateWithTimeIntervalSince1970:[credentials.expiration doubleValue]/
            1000];
        return nil;
    }];
}
```

# AWS Services can use caller credentials

[← Method Execution](#) /pets - POST - Integration Request

Delete Method


Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type ☒ Lambda Function


☐ HTTP Proxy

[Show advanced](#)

Lambda Region us-east-1 

Lambda Function APIGLambdaDemo 

Invoke with caller credentials ☒ 

Credentials cache Do not add caller credentials to cache key 

# Throttling and Caching

# API Throttling

Throttling helps you manage traffic to your backend

Throttle by developer-defined Requests/Sec limits

Requests over the limit are throttled

- HTTP 429 response

The generated SDKs retry throttled requests



# Caching of API Responses

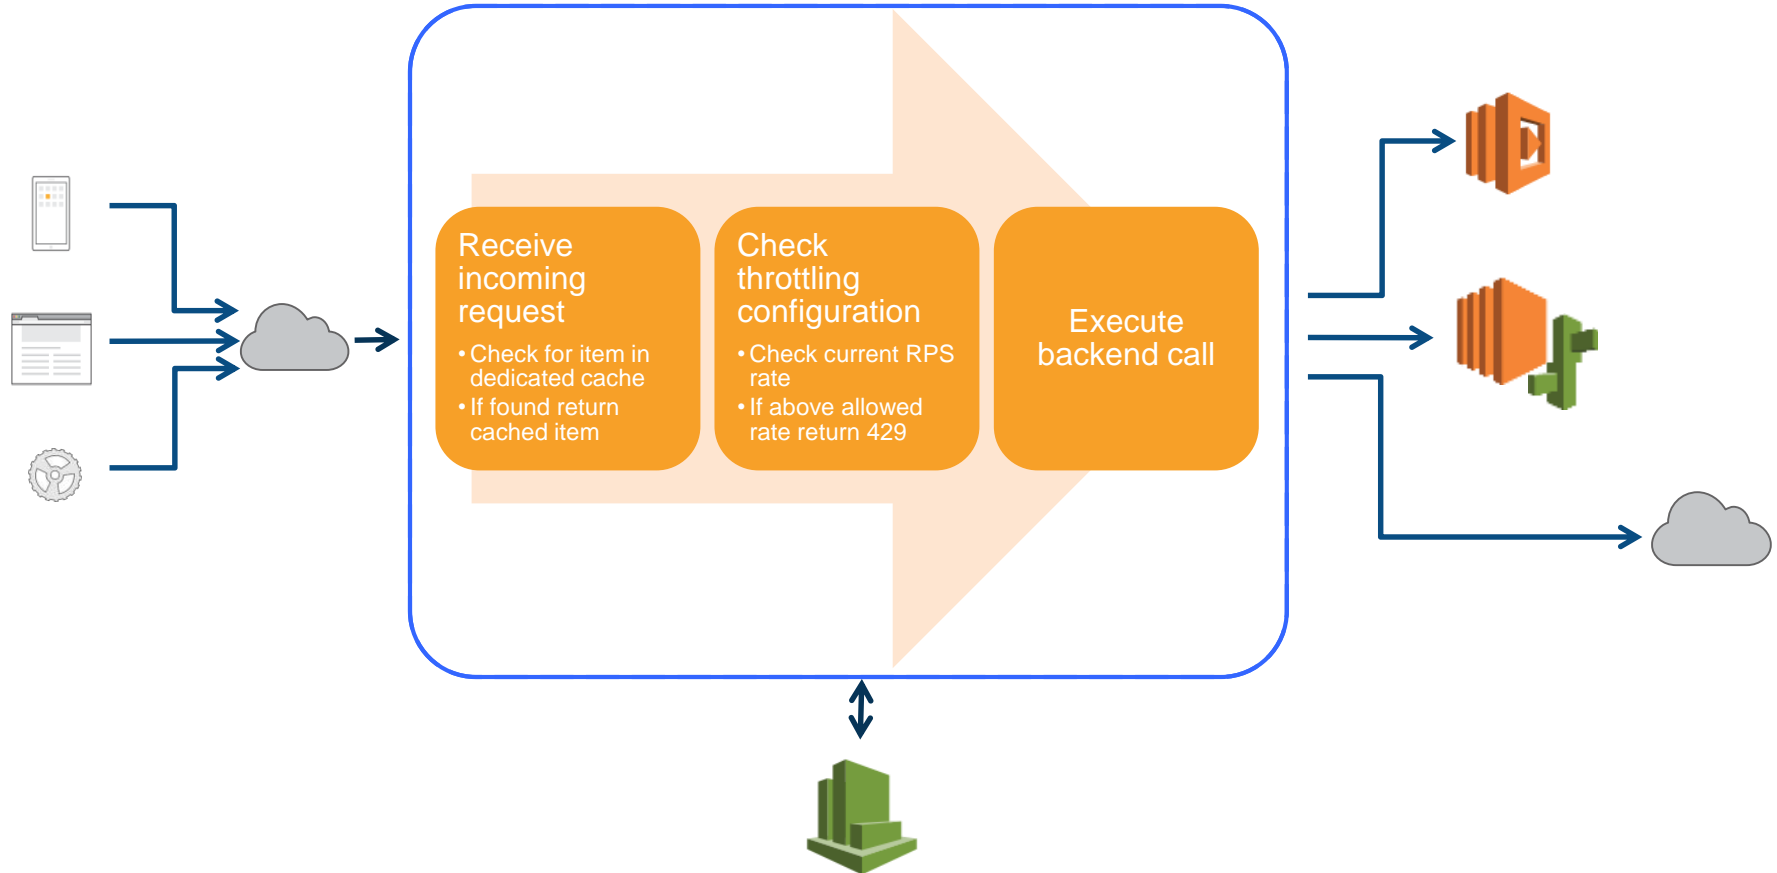
You can configure a cache key and the Time to Live (TTL) of the API response

Cached items are returned without calling the backend

A cache is dedicated to you, by stage

You can provision between 0.5GB to 237GB of cache

# Request processing workflow



# Input / Output Transformation

# Input / Output Transforms

Use Velocity Templates to transform data

Filter output results

- Remove private or unnecessary data
- Filter dataset size to improve API performance

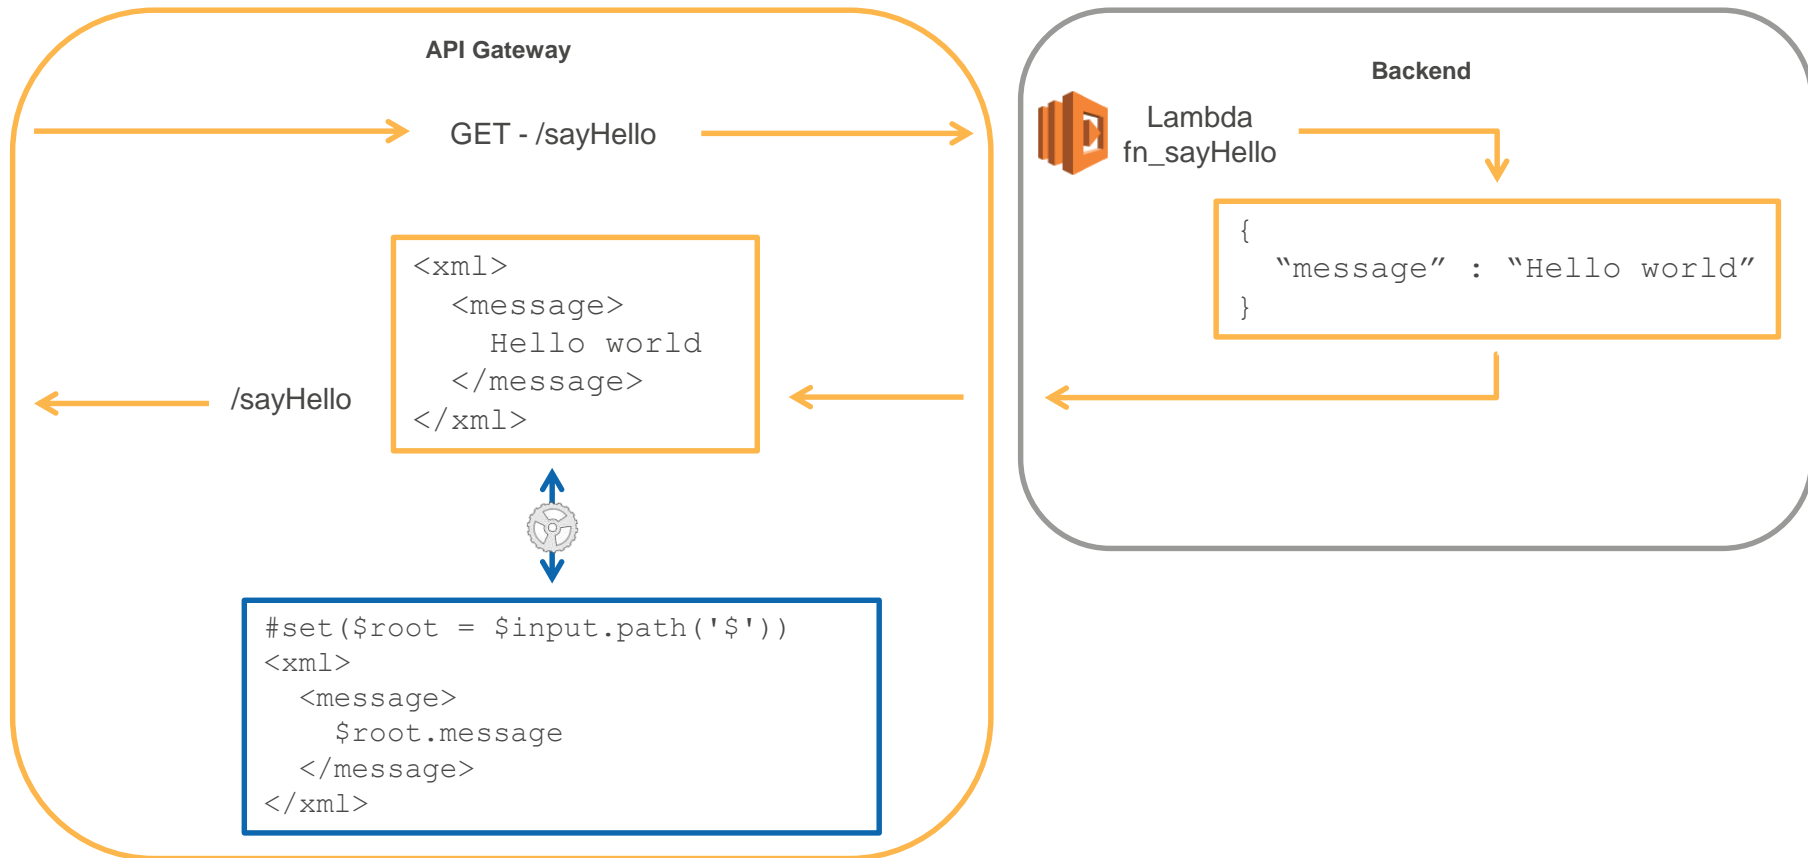
GET to POST

- Read all query string parameters from your GET request, and create a body to make a POST to your backend

JSON to XML

- Receive JSON input and transform it to XML for your backend
- Receive JSON from a Lambda function and transform it to XML

# Transform Example: JSON to XML



# For Loops and if Statements

```
1  #set($inputRoot = $input.path('$'))
2  {
3      "items" :[
4          #foreach($elem in $inputRoot.Items)
5              {
6                  "serviceName" : "$elem.serviceName.S",
7                  "dateCreated" : "$elem.dateCreated.N",
8                  "serviceId" : "$elem.serviceId.S"
9              }
10         #if($foreach.hasNext),#end
11     #end
12 ]
13 }
```

# One Template per Content/Type

## ▼ Mapping Templates

Content-Type	
application/json	⊖

[+ Add mapping template](#)

application/json

Mapping template 

Template ↗

```
1 {  
2   "action" : "com.amazonaws.apigatewaydemo.action.CreatePetD  
   emoAction",  
3   "body" : $input.json('$')  
4 }  
5
```

# SDK Generation



# API Models

Models are a JSON Schema representation of your API requests and responses

You can reuse models across multiple methods in your API

Models are used to generate objects for the client SDK

# Generate Client SDKs Based on Your APIs

SDKs are generated based on API deployments (Stages)

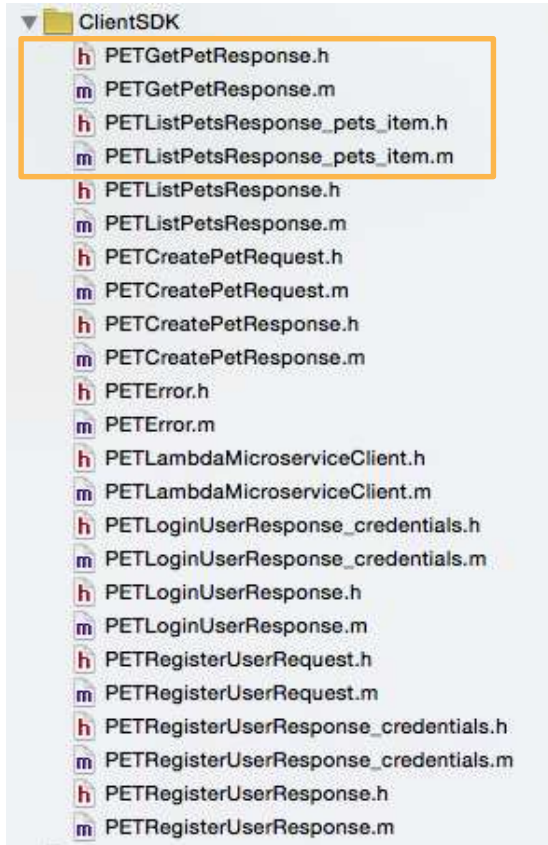
If Request and Response Models are defined, the SDK includes input and output marshalling of your methods

SDKs know how to handle throttling responses

SDKs also know how to sign requests with AWS temporary credentials (SigV4)

Support for Android, iOS, JavaScript, ...

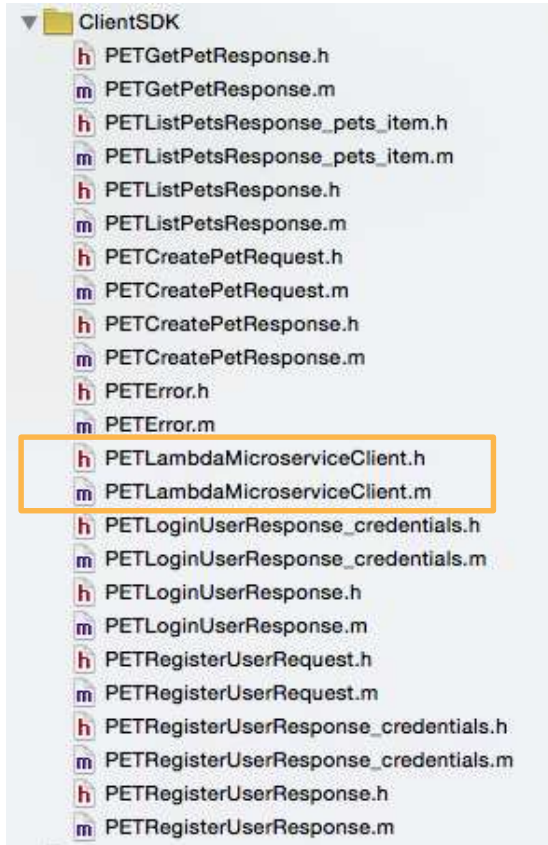
# Models are included in the SDK



Models are generated using their name

Nested items inherit the name of their parent model

# The Client SDK Declares All Methods



```
PETLambdaMicroserviceClient *client = [PETLambdaMicroserviceClient defaultClient];  
[[client petsGet] continueWithBlock:^(id(AWSTask *task) {  
    PETListPetsResponse *pets = task.result;  
    self.objects = [NSMutableArray arrayWithArray:pets.pets];  
    dispatch_async(dispatch_get_main_queue(), ^{  
        [self.tableView reloadData];  
        [hud hide:YES];  
    });  
    return nil;  
}];
```

# Thank You

Q&A